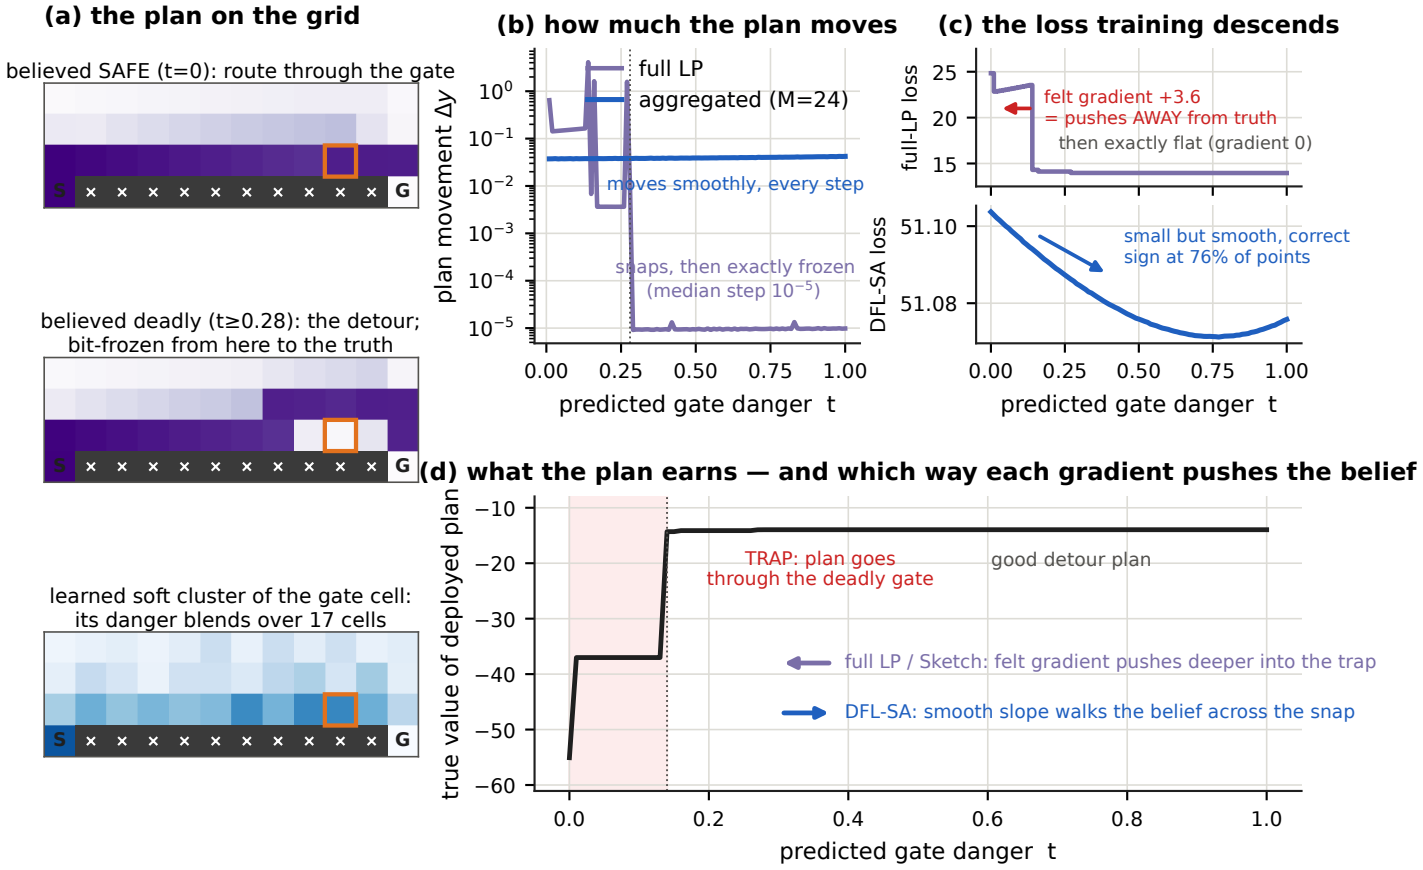
(a) the plan on the grid
believed SAFE (t=0): route through the gate
S
G
believed deadly (t≥0.28): the detour;
bit-frozen from here to the truth
S
G
learned soft cluster of the gate cell:
its danger blends over 17 cells
S
G
(b) how much the plan moves
plan movement Δy
full LP
aggregated (M=24)
moves smoothly, every step
snaps, then exactly frozen
(median step 10⁻⁵)
predicted gate danger t
(c) the loss training descends
full-LP loss
felt gradient +3.6
= pushes AWAY from truth
then exactly flat (gradient 0)
DFL-SA loss
small but smooth, correct
sign at 76% of points
predicted gate danger t
(d) what the plan earns — and which way each gradient pushes the belief
true value of deployed plan
TRAP: plan goes
through the deadly gate
good detour plan
full LP / Sketch: felt gradient pushes deeper into the trap
DFL-SA: smooth slope walks the belief across the snap
predicted gate danger t