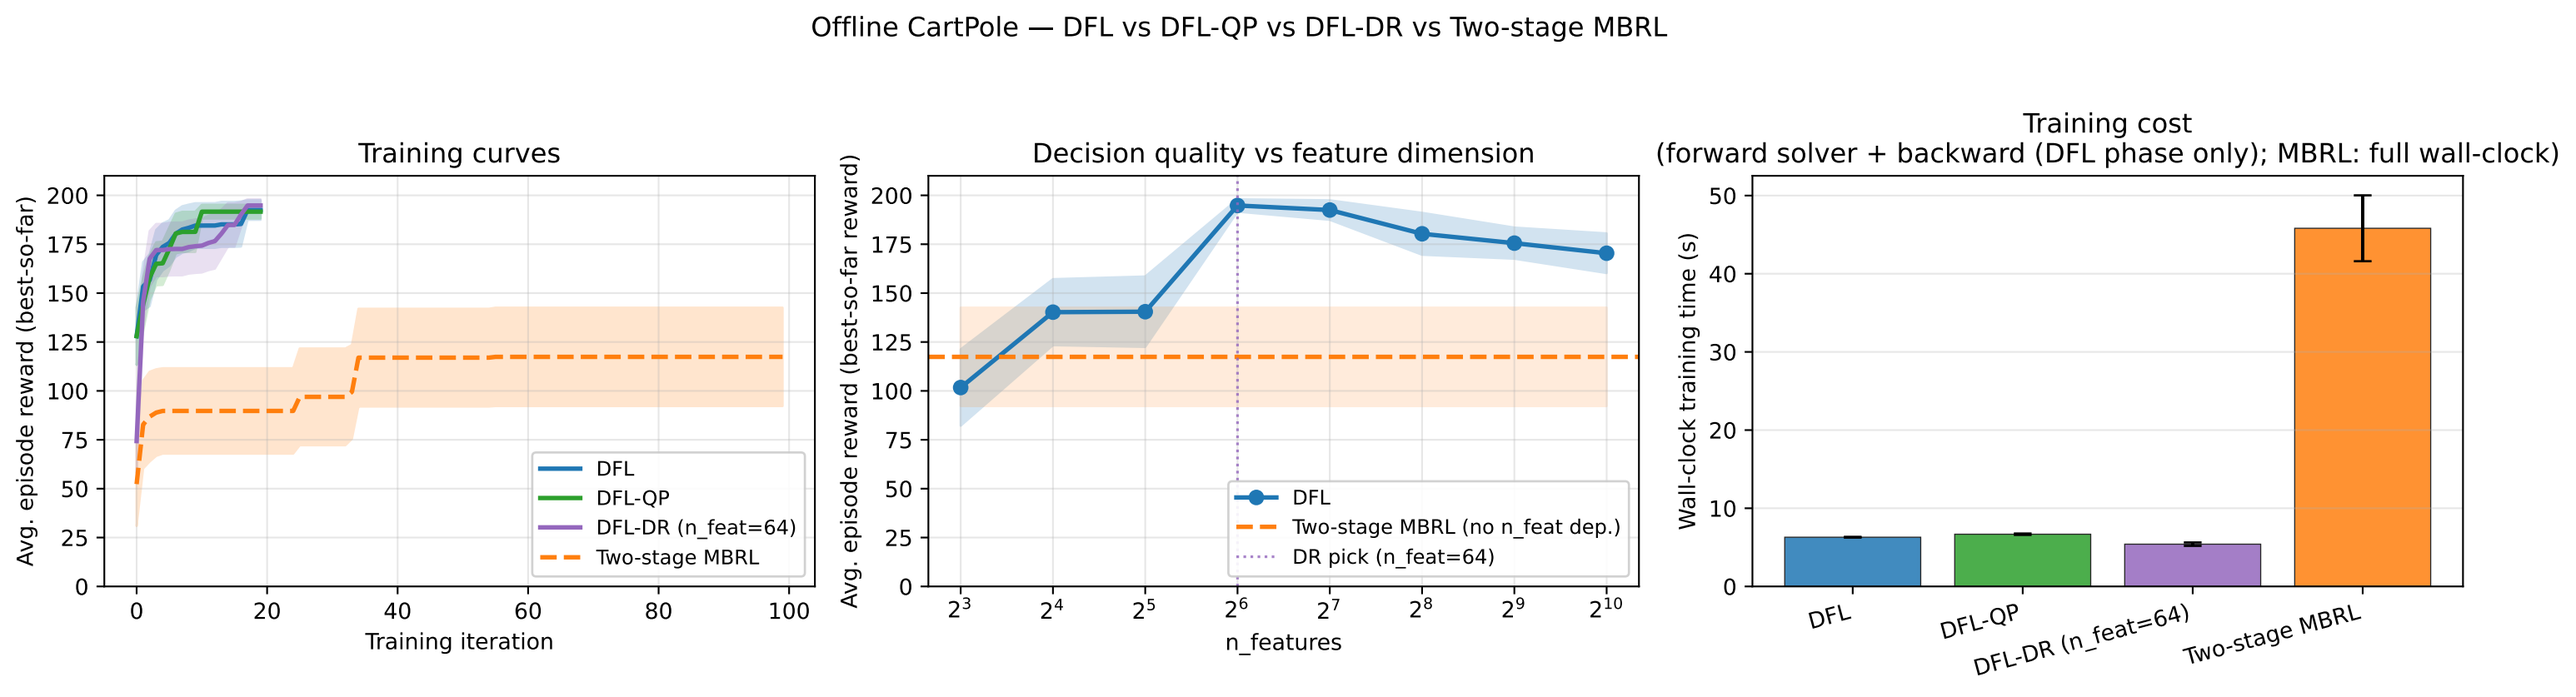
Offline CartPole — DFL vs DFL-QP vs DFL-DR vs Two-stage MBRL
Training curves
Avg. episode reward (best-so-far)
Training iteration
DFL
DFL-QP
DFL-DR (n_feat=64)
Two-stage MBRL
Decision quality vs feature dimension
Avg. episode reward (best-so-far reward)
n_features
DFL
Two-stage MBRL (no n_feat dep.)
DR pick (n_feat=64)
Training cost
(forward solver + backward (DFL phase only); MBRL: full wall-clock)
Wall-clock training time (s)
DFL
DFL-QP
DFL-DR (n_feat=64)
Two-stage MBRL