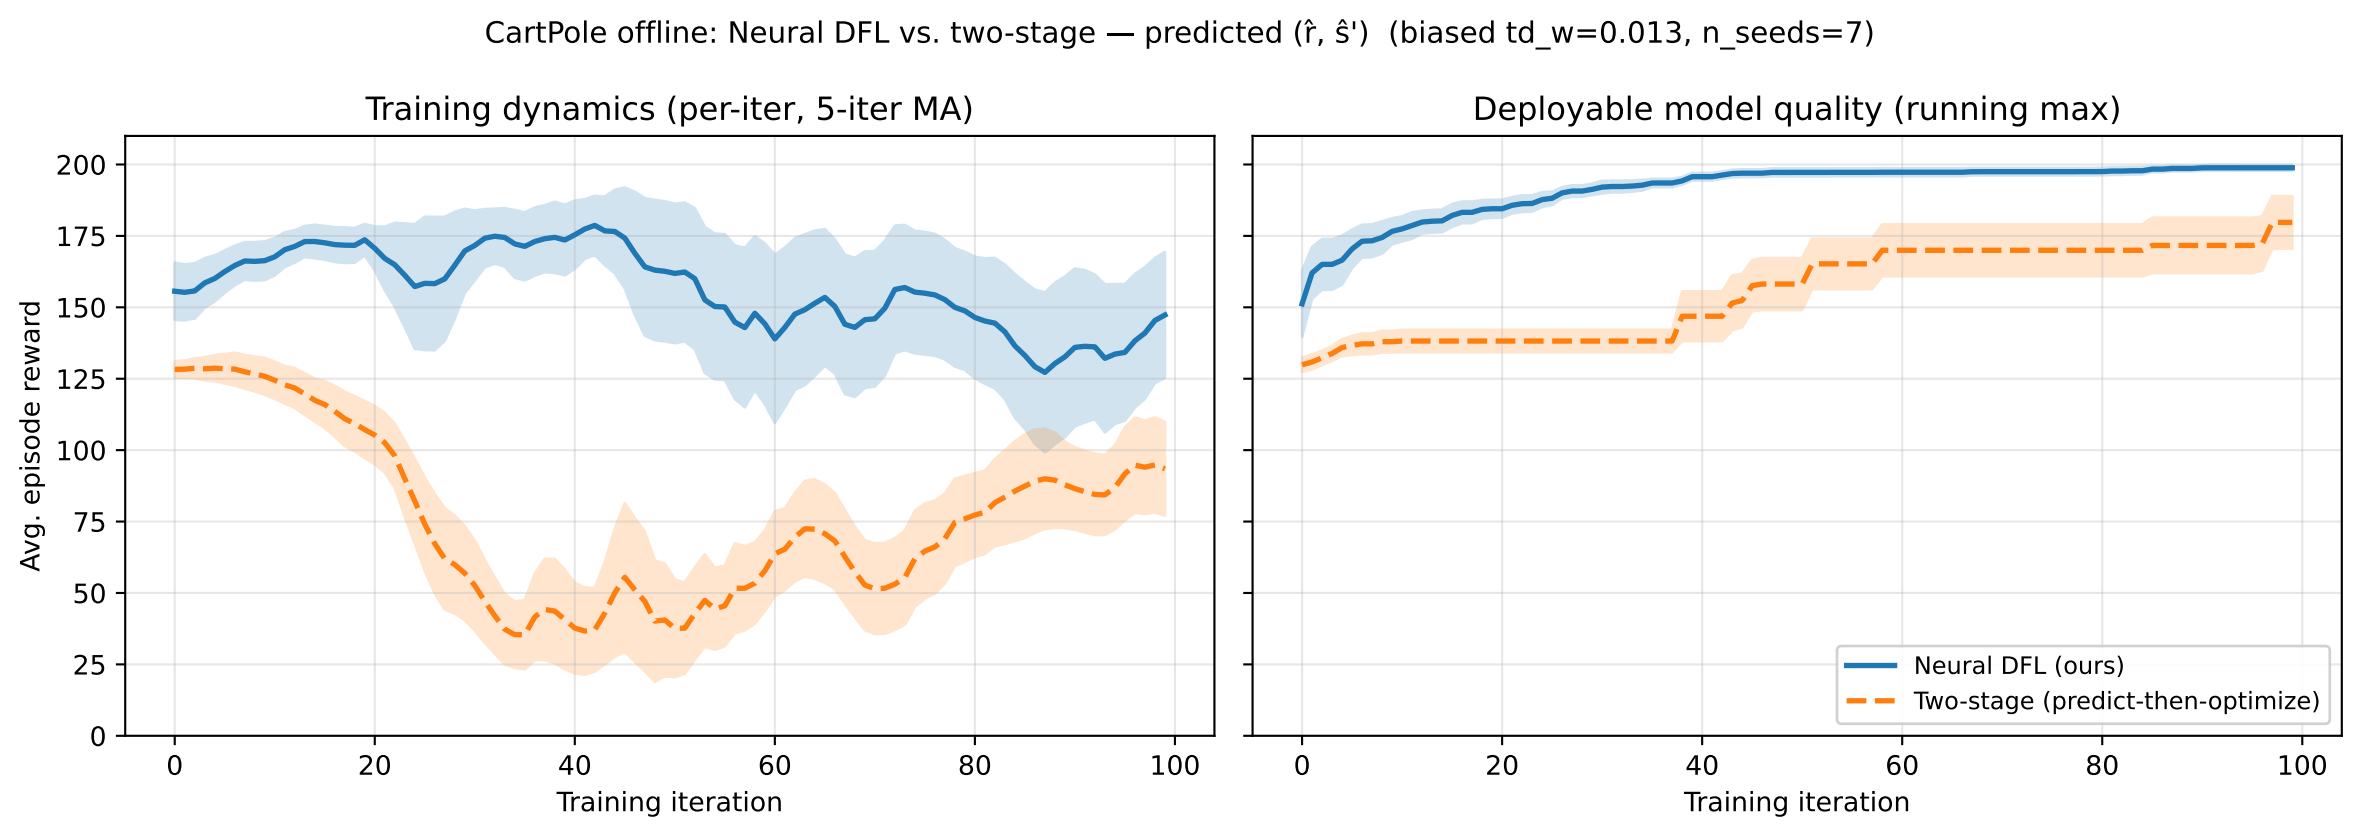
CartPole offline: Neural DFL vs. two-stage — predicted (r̂, ŝ') (biased td_w=0.013, n_seeds=7)
Training dynamics (per-iter, 5-iter MA)
Deployable model quality (running max)
Avg. episode reward
Training iteration
Training iteration
Neural DFL (ours)
Two-stage (predict-then-optimize)