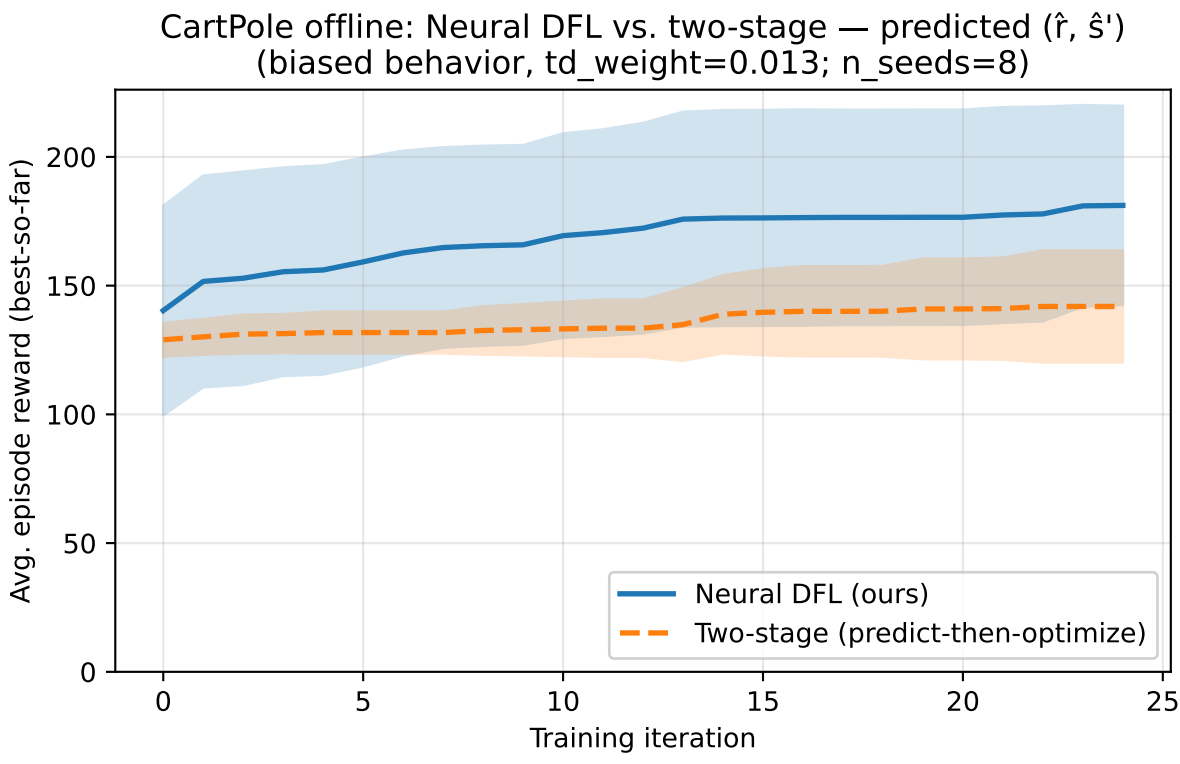
CartPole offline: Neural DFL vs. two-stage — predicted (r̂, ŝ')
(biased behavior, td_weight=0.013; n_seeds=8)
Avg. episode reward (best-so-far)
Training iteration
0
50
100
150
200
0
5
10
15
20
25
Neural DFL (ours)
Two-stage (predict-then-optimize)