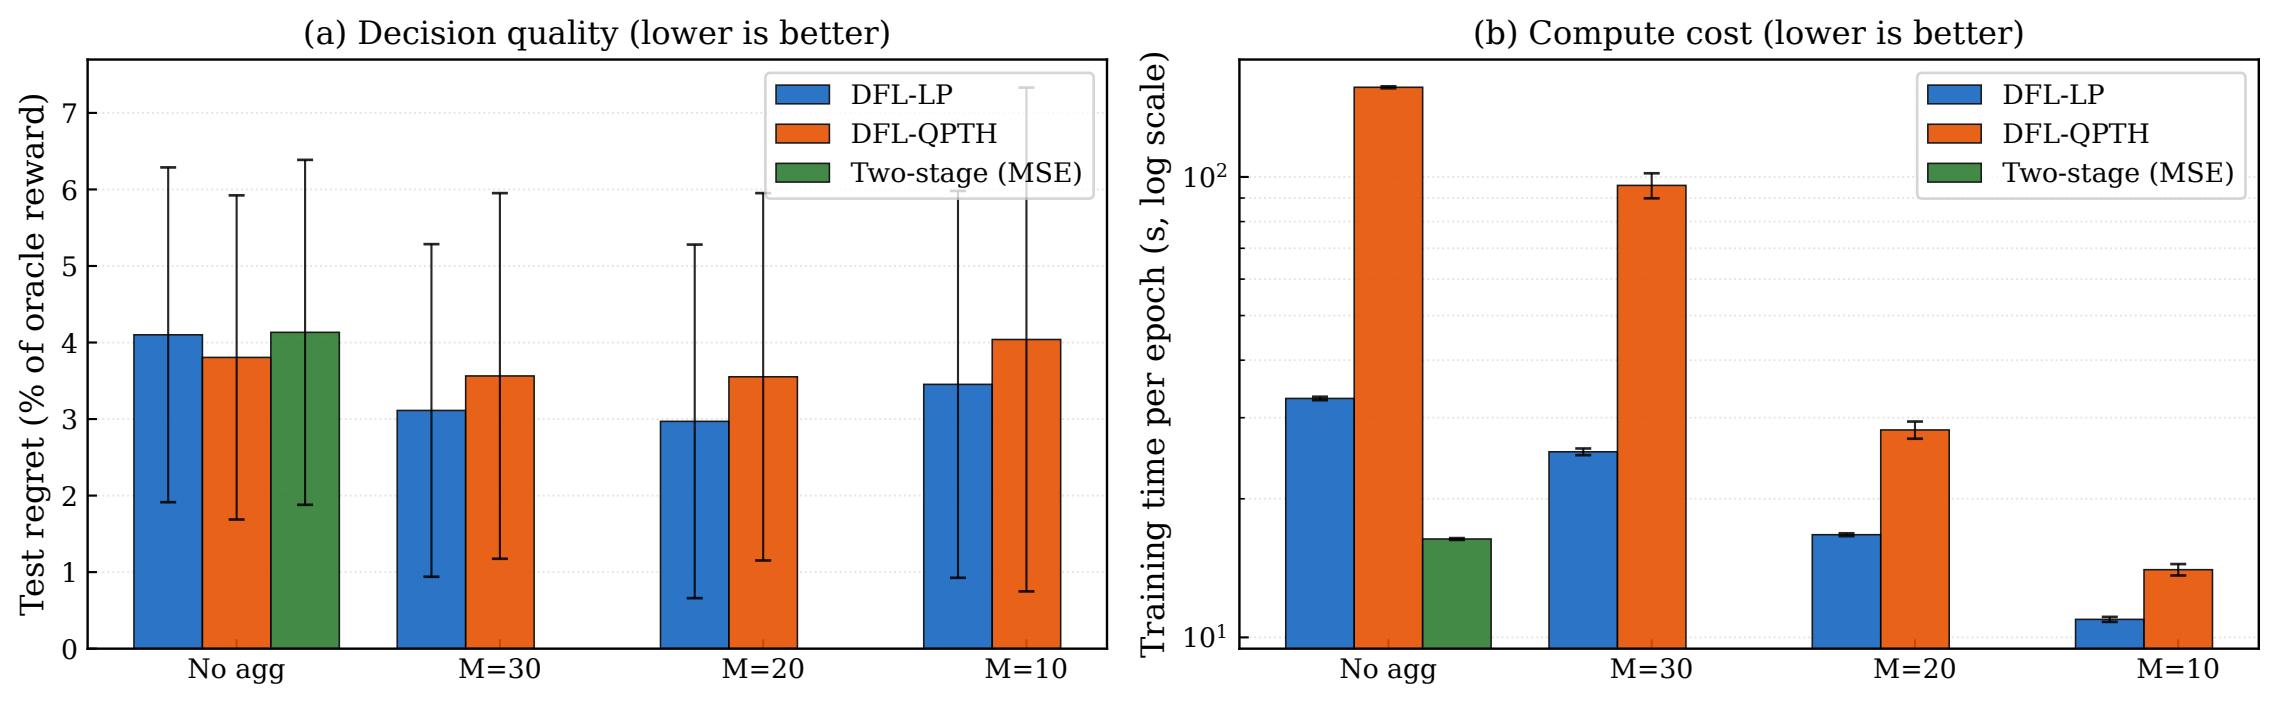
(a) Decision quality (lower is better)
Test regret (% of oracle reward)
0
1
2
3
4
5
6
7
DFL-LP
DFL-QPTH
Two-stage (MSE)
No agg
M=30
M=20
M=10
(b) Compute cost (lower is better)
Training time per epoch (s, log scale)
10^2
10^1
DFL-LP
DFL-QPTH
Two-stage (MSE)
No agg
M=30
M=20
M=10