DFL Objective: True Reward Under Predicted Policy
(Gray lines = LP basis changes; peak should be near λ_true)

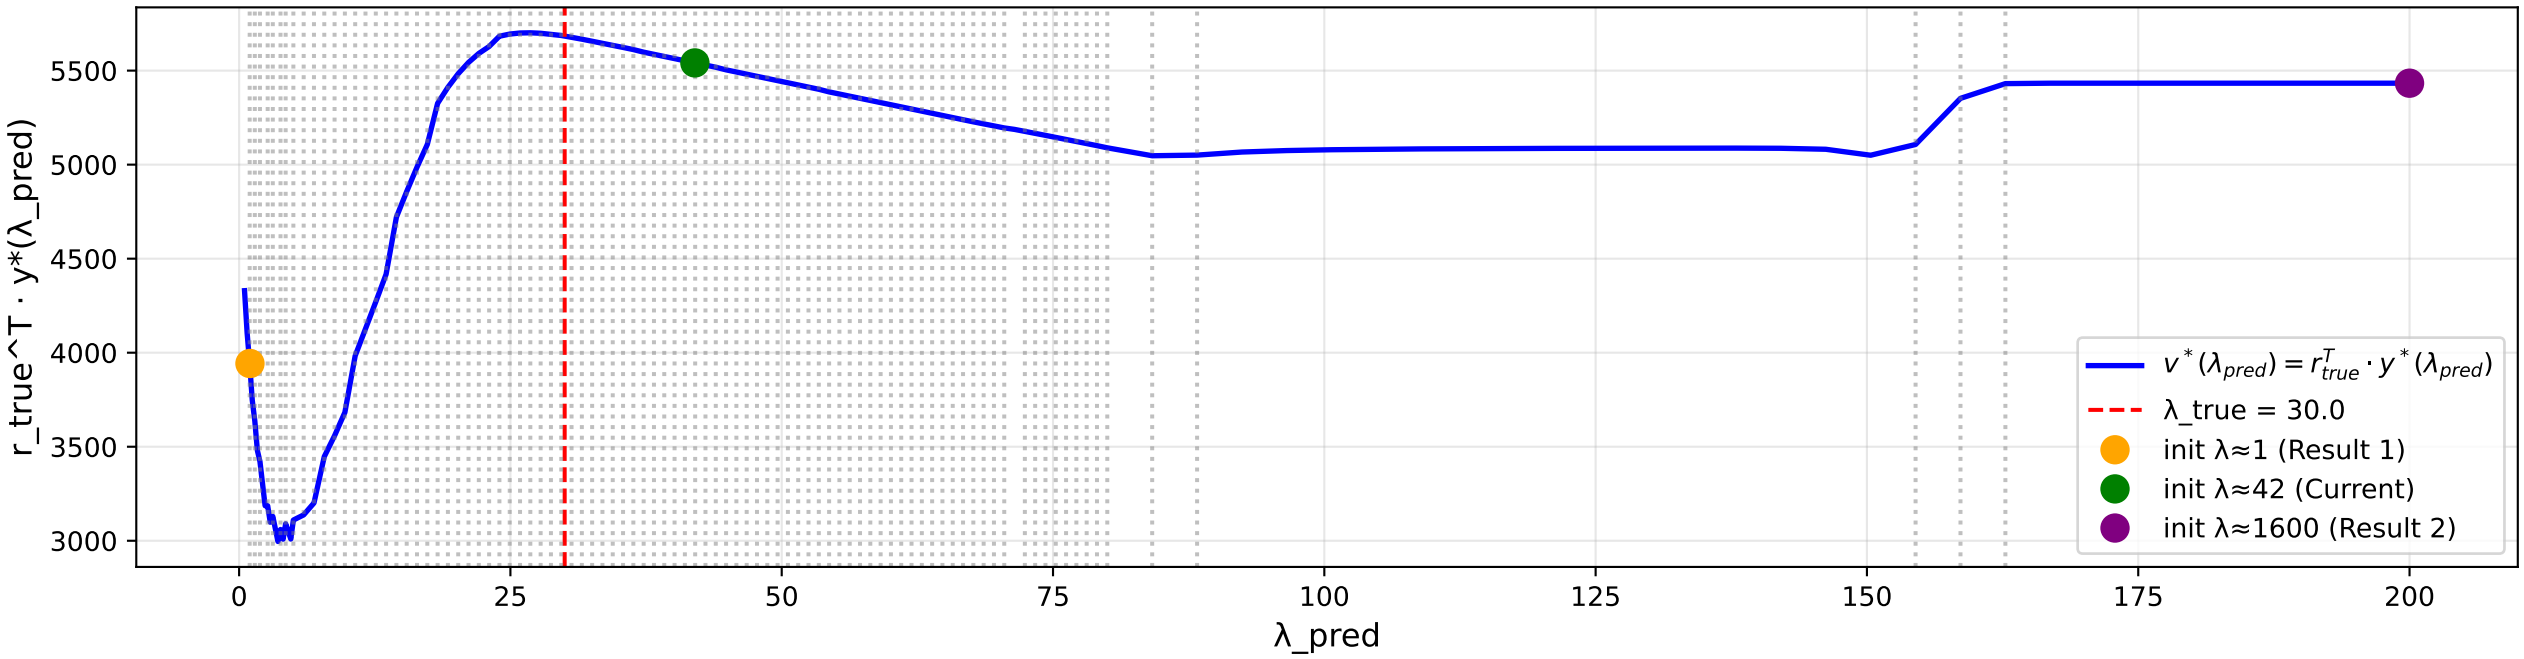

DFL Gradient: dL/dλ
(Negative → SGD increases λ; Positive → SGD decreases λ)

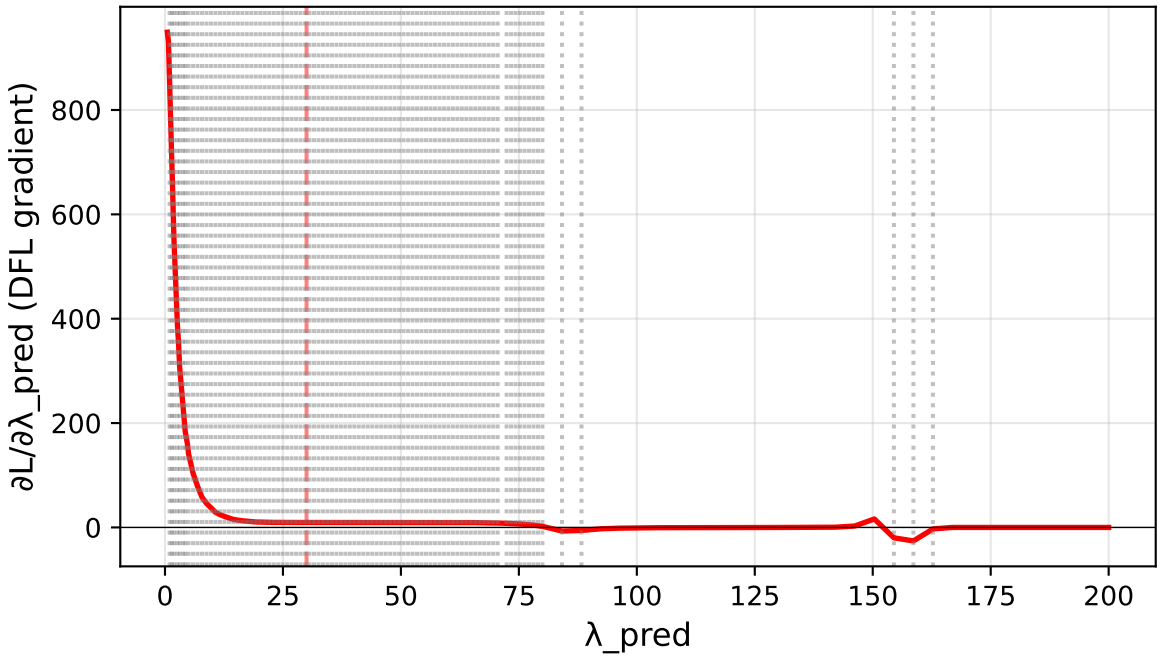

Predicted LP Objective
(what the LP maximizes internally)

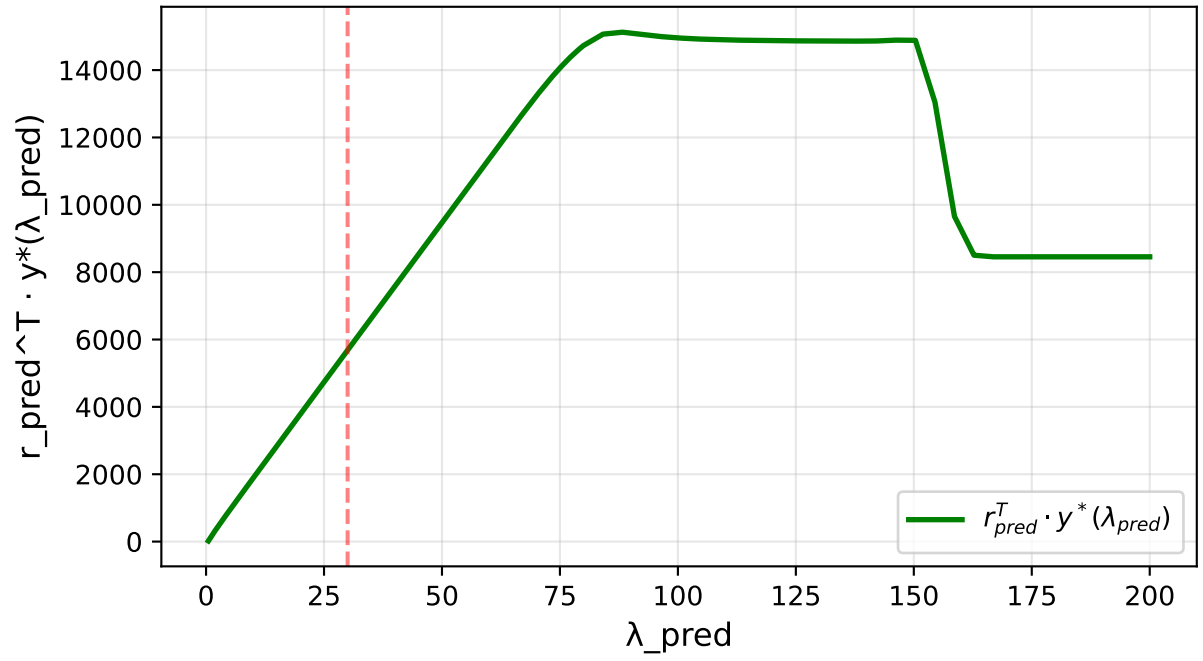
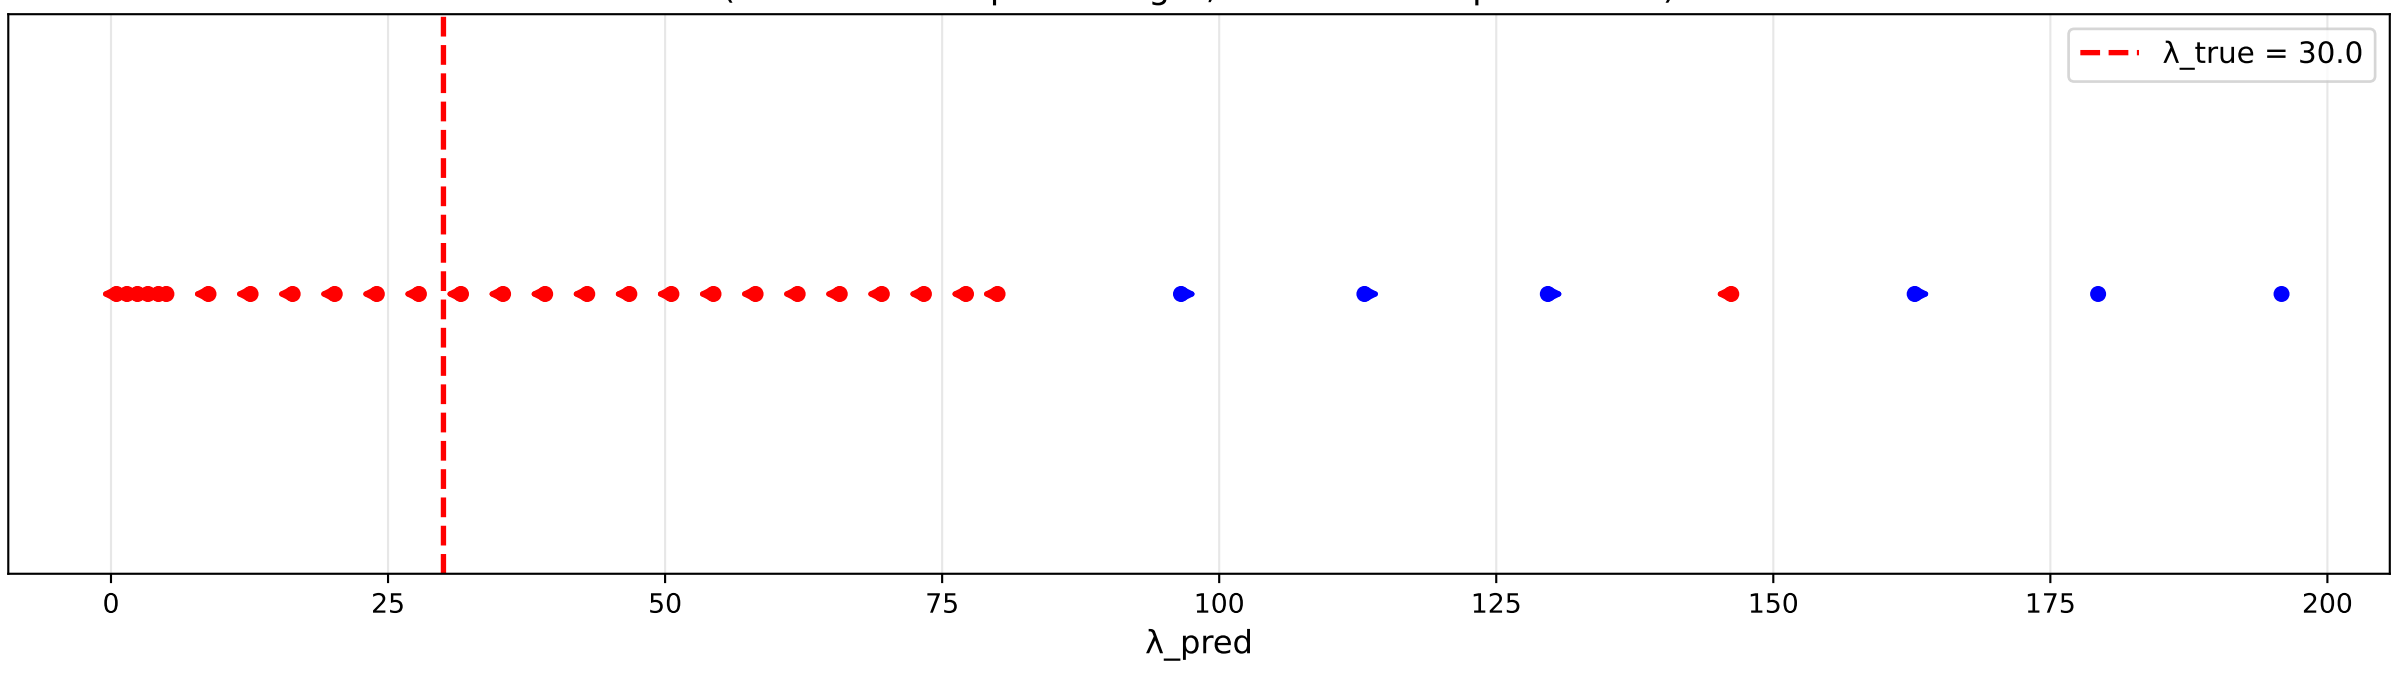
SGD Update Direction at Each λ_pred
(Blue → = SGD pushes right; Red ← = SGD pushes left)